

# **SENTIMENT ANALYSIS OF AMAZON PRODUCT REVIEWS FOR CONSUMER INSIGHTS**

---

**DSCI 6004-2: Natural Language Processing**  
Khaled Sayed, PhD – Assistant Professor

---

**TEAM:**

**MADHAVI KANCHAM**

**VINAYKUMAR REDDY MOKU**

**NITHISH KUMAR MAYAVAN**

A solid orange horizontal bar at the bottom of the slide.

# Project Overview

---

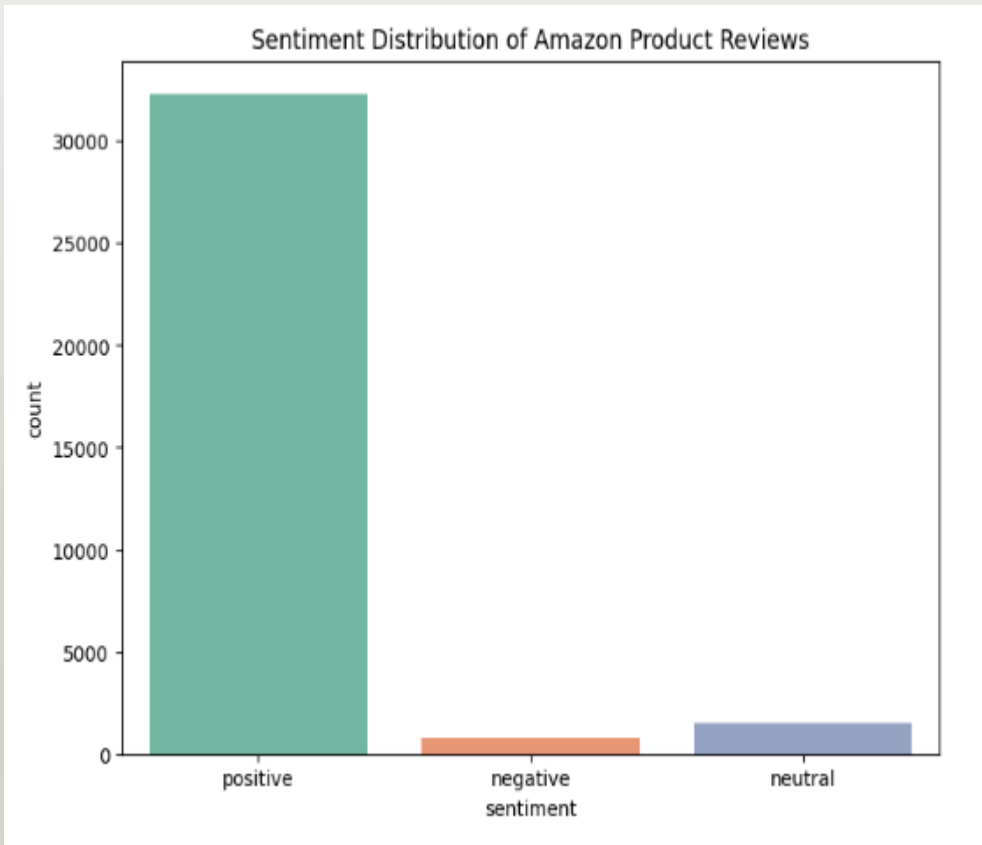
- Classify Amazon product reviews as positive, neutral or negative using sentiment analysis.
- Predict consumer satisfaction and identify trends in reviews to provide insights for product improvement and marketing strategies.
- Leverage NLP and machine learning algorithms to analyze largescale review data and automate sentiment classification.

# Dataset


The dataset for this study was collected from Amazon through Kaggle with over 34000 rows and 21 variables on reviews. Available at:  
<https://www.kaggle.com/datasets/bittlingmayer/amazonreviews>

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34660 entries, 0 to 34659
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    34660 non-null  object
1   name                                27900 non-null  object
2   asins                               34658 non-null  object
3   brand                               34660 non-null  object
4   categories                           34660 non-null  object
5   keys                                 34660 non-null  object
6   manufacturer                         34660 non-null  object
7   reviews.date                        34621 non-null  object
8   reviews.dateAdded                   24039 non-null  object
9   reviews.dateSeen                    34660 non-null  object
10  reviews.didPurchase                 1 non-null      object
11  reviews.doRecommend                 34066 non-null  object
12  reviews.id                          1 non-null      float64
13  reviews.numHelpful                  34131 non-null  float64
14  reviews.rating                      34627 non-null  float64
15  reviews.sourceURLs                  34660 non-null  object
16  reviews.text                        34659 non-null  object
17  reviews.title                       34654 non-null  object
18  reviews.userCity                    0 non-null      float64
19  reviews.userProvince                0 non-null      float64
20  reviews.username                    34653 non-null  object
dtypes: float64(5), object(16)
memory usage: 5.6+ MB
```

# Sentiment Distribution




# Most Frequent Words



# Distribution of Ratings

---



# Preprocessing

---

- Tokenization and Vectorization: Used `TfidfVectorizer(max_features=10000)` to transform preprocessed reviews into TF-IDF feature vectors.
- Sentiment Label Encoding: Converted sentiment labels (negative, neutral, positive) into numeric values (0, 1, 2).
- Data Splitting: Split the dataset into training (80%) and testing (20%) sets using `train_test_split` with stratified sampling.

```
Training data shape: (27700, 10000)  
Test data shape: (6926, 10000)
```

# Feedforward Neural Network (Multilayer Perceptron)

---

- The SentimentNN class defines a simple feedforward neural network with one hidden layer (fc1), ReLU activation, and an output layer (fc2) followed by a softmax activation for multi-class classification.
- It processes the input data (features) to predict sentiment labels (negative, neutral, positive).
- Loss and Optimization: Uses CrossEntropyLoss for multi-class classification and Adam optimizer for efficient training.
- Training Setup: Input dimension matches feature count, 128 hidden units, 3 output classes (negative, neutral, positive), with a learning rate of 0.001

```
14]: class SentimentNN(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(SentimentNN, self).__init__()
        self.fc1 = nn.Linear(input_dim, hidden_dim)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_dim, output_dim)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.softmax(x)
        return x

# Define the model, loss function, and optimizer
input_dim = X_train.shape[1] # Number of features
hidden_dim = 128 # Number of hidden units
output_dim = 3 # Three classes: negative, neutral, positive

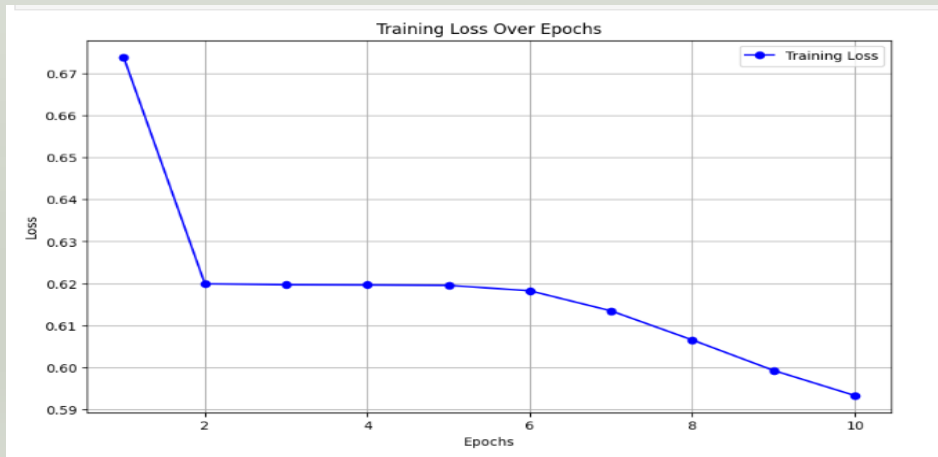
model = SentimentNN(input_dim, hidden_dim, output_dim)

# Set up Loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```



# Neural Network Performance

- The model shows a steady decrease in loss from 0.6738 to 0.5934 over 10 epochs, indicating improvement during training.
- Despite the low loss, performance is skewed, with the model achieving 93.17% accuracy overall, excelling in positive sentiment prediction but struggling with negative (45% precision, 27% recall) and neutral (30% precision, 11% recall) classes.



Test Accuracy: 0.9317

	precision	recall	f1-score	support
negative	0.45	0.27	0.33	162
neutral	0.30	0.11	0.16	300
positive	0.95	0.99	0.97	6464
accuracy			0.93	6926
macro avg	0.57	0.45	0.49	6926
weighted avg	0.91	0.93	0.92	6926

# Neural Network deployment

---

Running on local URL: <http://127.0.0.1:7870>

To create a public link, set ``share=True`` in ``launch()``.

## Amazon Product Review Sentiment Analysis

Enter an Amazon product review and get the predicted sentiment (negative, neutral, or positive).

Enter your review here

This product is excellent! It works as described, has great quality, and exceeded my expectations. Highly recommended!

Clear

Submit

Predicted Sentiment

Predicted sentiment for the review: positive

Flag

# LSTM MODEL

---

```
class LSTMModel(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, n_layers, dropout_prob):
        super(LSTMModel, self).__init__()

        self.embedding = nn.Embedding(vocab_size, embedding_dim)

        self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers, batch_first=True, dropout=dropout_prob)

        self.fc = nn.Linear(hidden_dim, output_dim)

        # Dropout Layer to prevent overfitting
        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, (hidden, cell) = self.lstm(embedded)
        # We use the last hidden state as the representation for the entire sequence
        hidden = hidden[-1, :, :]
        out = self.fc(self.dropout(hidden))
        return out
```

The LSTMModel defines an embedding layer, an LSTM layer, a fully connected layer, and a dropout layer to process sequential input and output sentiment predictions.

# LTSM MODEL Performance

---

The model achieves 95% accuracy, performing excellently on positive sentiment (95% precision, 100% recall), but poorly on negative and neutral sentiments, resulting in low precision, recall, and F1-scores for those classes.

	precision	recall	f1-score	support
negative	0.00	0.00	0.00	149
neutral	0.10	0.00	0.01	222
positive	0.95	1.00	0.97	6555
accuracy			0.95	6926
macro avg	0.35	0.33	0.33	6926
weighted avg	0.90	0.95	0.92	6926

# LTSM Deployment

---

Running on local URL: `http://127.0.0.1:7870`

To create a public link, set ``share=True`` in ``launch()``.

## Amazon Product Review Sentiment Analysis

Enter an Amazon product review and get the predicted sentiment (negative, neutral, or positive).

Enter your review here

This product is disappointing. It broke quickly, doesn't work as expected, and is of poor quality. Not recommended.

Clear

Submit

Predicted Sentiment

Predicted sentiment for the review: negative

Flag

# DistilBERT Model

---

The model loads a pretrained DistilBERT model and tokenizer for sequence classification, moving the model to the appropriate device (GPU/CPU) for efficient processing

```
# Load the pretrained DistilBERT model and tokenizer
model = DistilBertForSequenceClassification.from_pretrained('distilbert-base-uncased', num_labels=3)
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
```

# DistilBERT deployment

---

Running on local URL: <http://127.0.0.1:7873>

To create a public link, set ``share=True`` in ``launch()``.

## Sentiment Analysis of Product Reviews

Enter a product review, and the model will predict whether the sentiment is 'negative', 'neutral', or 'positive'.

Enter Review

The product is okay. It works as expected but doesn't stand out in terms of quality or features

Clear

Predicted Sentiment

negative



Flag

# BERT Model for sentimental

The BertForSequenceClassification model consists of a BertModel with embeddings, encoder layers, and a pooler, followed by a dropout layer and a classifier that outputs predictions for three classes.

The encoder uses 12 transformer layers, each with self-attention and feed-forward components, and the pooler produces the final hidden state for classification.

```
[30]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
          (intermediate): BertIntermediate(
            (dense): Linear(in_features=768, out_features=3072, bias=True)
            (intermediate_act_fn): GELUActivation()
          )
          (output): BertOutput(
            (dense): Linear(in_features=3072, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
    (dropout): Dropout(p=0.1, inplace=False)
    (classifier): Linear(in_features=768, out_features=3, bias=True)
  )
)
```



# BERT Model for sentimental

---

Running on local URL: <http://127.0.0.1:7874>

To create a public link, set ``share=True`` in ``launch()``.

## Sentiment Analysis of Product Reviews

Enter a product review, and the model will predict whether the sentiment is 'negative', 'neutral', or 'positive'.

Enter Review

This product is disappointing. It stopped working after a few uses and doesn't meet the advertised expectations. Not recommended

Clear

Predicted Sentiment

negative

Flag

---

**THANK YOU**

