

SENTIMENT ANALYSIS OF AMAZON PRODUCT REVIEWS FOR CONSUMER INSIGHTS

```
In [1]: # Import necessary Libraries
import pandas as pd
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt
import seaborn as sns
from nltk.corpus import stopwords
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report
from torch.utils.data import DataLoader, TensorDataset
from torch.nn.utils.rnn import pad_sequence
from transformers import BertTokenizer
from wordcloud import WordCloud
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
# Download necessary NLTK resources
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('common_words')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\MAPILI\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\MAPILI\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Error loading common_words: Package 'common_words' not
[nltk_data] found in index
```

Out[1]: False

Data Preprocessing

```
In [2]: import pandas as pd
import warnings
warnings.filterwarnings("ignore")
# Load the dataset
df = pd.read_csv('Amazon_Reviews.csv')

# Display the first few rows of the dataset
df.head()
```

Out[2]:

		id	name	asins	brand	categories	keys	manufacturer	reviews.date	review
0	AVqklhwDv8e3D1O-lebb		All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	Amazon	2017-01-13T00:00:00.000Z	
1	AVqklhwDv8e3D1O-lebb		All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	Amazon	2017-01-13T00:00:00.000Z	
2	AVqklhwDv8e3D1O-lebb		All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	Amazon	2017-01-13T00:00:00.000Z	
3	AVqklhwDv8e3D1O-lebb		All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	Amazon	2017-01-13T00:00:00.000Z	
4	AVqklhwDv8e3D1O-lebb		All-New Fire HD 8 Tablet, 8 HD Display, Wi-Fi...	B01AHB9CN2	Amazon	Electronics,iPad & Tablets,All Tablets,Fire Ta...	841667104676,amazon/53004484,amazon/b01ahb9cn2...	Amazon	2017-01-12T00:00:00.000Z	

5 rows × 21 columns

In [3]:

```
# Display summary information about the dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34660 entries, 0 to 34659
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     34660 non-null  object
1   name                   27900 non-null  object
2   asins                   34658 non-null  object
3   brand                   34660 non-null  object
4   categories              34660 non-null  object
5   keys                    34660 non-null  object
6   manufacturer            34660 non-null  object
7   reviews.date            34621 non-null  object
8   reviews.dateAdded       24039 non-null  object
9   reviews.dateSeen        34660 non-null  object
10  reviews.didPurchase      1 non-null      object
11  reviews.doRecommend     34066 non-null  object
12  reviews.id              1 non-null      float64
13  reviews.numHelpful       34131 non-null  float64
14  reviews.rating           34627 non-null  float64
15  reviews.sourceURLs       34660 non-null  object
16  reviews.text             34659 non-null  object
17  reviews.title            34654 non-null  object
18  reviews.userCity         0 non-null      float64
19  reviews.userProvince    0 non-null      float64
20  reviews.username         34653 non-null  object
dtypes: float64(5), object(16)
memory usage: 5.6+ MB
```

Explanatory Data Analysis (EDA)

In [4]:

```
# Handle missing values and drop rows with missing reviews or ratings
df = df.dropna(subset=['reviews.text', 'reviews.rating'])

# Map ratings to sentiment labels (negative, neutral, positive)
df['sentiment'] = df['reviews.rating'].apply(lambda x: 'negative' if x <= 2 else 'neutral' if x == 3 else 'positive')

# Display sentiment distribution
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='sentiment', palette='Set2')
plt.title('Sentiment Distribution of Amazon Product Reviews')
plt.show()
```

A bar chart titled 'sentiment' showing the count of three sentiment categories. The y-axis is labeled 'count' and ranges from 0 to 30,000 in increments of 5,000. The x-axis lists the sentiment categories: 'positive', 'negative', and 'neutral'. The 'positive' bar is teal and reaches approximately 32,000. The 'negative' bar is orange and reaches approximately 1,000. The 'neutral' bar is blue and reaches approximately 1,500.

sentiment	count
positive	32000
negative	1000
neutral	1500

[illegible]

```
In [6]: # Sentiment-Specific Word Cloud (Positive, Negative, Neutral)
# Positive Reviews
positive_reviews = df[df['sentiment'] == 'positive']
positive_wordcloud = WordCloud(width=800, height=400, background_color='white', stopwords=stop_words, max_words=200).generate(' '.join(positive_r

plt.figure(figsize=(10, 6))
plt.imshow(positive_wordcloud, interpolation='bilinear')
plt.title('Word Cloud for Positive Reviews')
plt.axis('off')
plt.show()
```


[illegible][illegible]

```
In [9]: # Frequency Distribution of words
from nltk.probability import FreqDist


all_words = ' '.join(df['processed_reviews']).split()
fdist = FreqDist(all_words)

# Display top 20 most frequent words
top_words = fdist.most_common(20)
top_words_df = pd.DataFrame(top_words, columns=['Word', 'Frequency'])

plt.figure(figsize=(12, 6))
sns.barplot(x='Frequency', y='Word', data=top_words_df, palette='Set2')
plt.title('Top 20 Most Frequent Words')
plt.show()
```

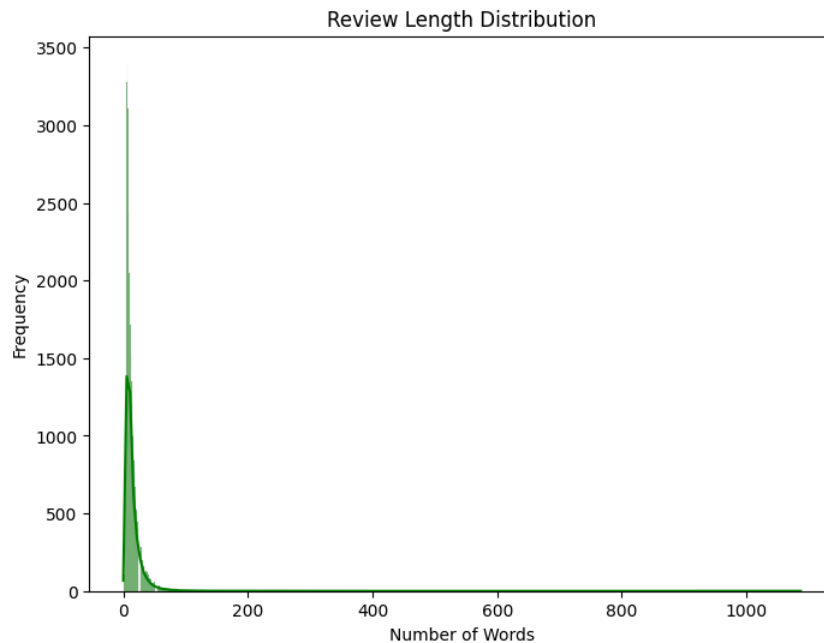


```
In [10]: # EDA: Histogram for Rating Distribution
plt.figure(figsize=(8, 6))
sns.histplot(df['reviews.rating'], kde=False, bins=5, color='skyblue')
plt.title('Distribution of Ratings')
plt.xlabel('Rating')
plt.ylabel('Frequency')
plt.show()
```



```
In [11]: # EDA: Review Length Distribution (number of words)
df['review_length'] = df['processed_reviews'].apply(lambda x: len(x.split()))

plt.figure(figsize=(8, 6))
sns.histplot(df['review_length'], kde=True, color='green')
plt.title('Review Length Distribution')
plt.xlabel('Number of Words')
plt.ylabel('Frequency')
plt.show()
```



TfidfVectorizer for text vectorization

```
In [12]: # Initialize TfidfVectorizer for text vectorization
vectorizer = TfidfVectorizer(max_features=10000)

# Fit the vectorizer and transform the text data
X = vectorizer.fit_transform(df['processed_reviews']).toarray()

# Convert sentiment labels to numeric values (negative: 0, neutral: 1, positive: 2)
y = df['sentiment'].map({'negative': 0, 'neutral': 1, 'positive': 2}).values
```

Split the data

```
In [13]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)

# Display the shape of train and test sets
print(f"Training data shape: {X_train.shape}")
print(f"Test data shape: {X_test.shape}")
```

Training data shape: (27700, 10000)
Test data shape: (6926, 10000)

In []:

Model 1: Simple Feedforward Neural Network (Multilayer Perceptron, or MLP) for Sentiment Analysis

```
In [14]: class SentimentNN(nn.Module):
    def __init__(self, input_dim, hidden_dim, output_dim):
        super(SentimentNN, self).__init__()
        self.fc1 = nn.Linear(input_dim, hidden_dim)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_dim, output_dim)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.softmax(x)
        return x

# Define the model, Loss function, and optimizer
input_dim = X_train.shape[1] # Number of features
hidden_dim = 128 # Number of hidden units
output_dim = 3 # Three classes: negative, neutral, positive

model = SentimentNN(input_dim, hidden_dim, output_dim)

# Set up Loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
In [1]: import torch
import torch.nn as nn
from torchsummary import summary

# Define the SentimentNN class
class SentimentNN(nn.Module):
```

```

def __init__(self, input_dim, hidden_dim, output_dim):
    super(SentimentNN, self).__init__()
    self.fc1 = nn.Linear(input_dim, hidden_dim)
    self.relu = nn.ReLU()
    self.fc2 = nn.Linear(hidden_dim, output_dim)
    self.softmax = nn.Softmax(dim=1)

def forward(self, x):
    x = self.fc1(x)
    x = self.relu(x)
    x = self.fc2(x)
    x = self.softmax(x)
    return x

# Example dimensions (replace with actual X_train.shape[1] if known)
input_dim = 10000 # Assuming 10,000 features
hidden_dim = 128 # Hidden Layer units
output_dim = 3 # Number of sentiment classes

# Initialize the model
model = SentimentNN(input_dim, hidden_dim, output_dim)

# Print the model summary
summary(model, input_size=(1, input_dim))

```

```

-----
Layer (type)                Output Shape          Param #
-----
      Linear-1                [-1, 1, 128]          1,280,128
      ReLU-2                  [-1, 1, 128]           0
      Linear-3                [-1, 1, 3]             387
      Softmax-4               [-1, 1, 3]             0
-----
Total params: 1,280,515
Trainable params: 1,280,515
Non-trainable params: 0
-----
Input size (MB): 0.04
Forward/backward pass size (MB): 0.00
Params size (MB): 4.88
Estimated Total Size (MB): 4.92
-----

```

```

In [15]: # Initialize lists to store loss values
train_losses = []

# Convert data to torch tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.long)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)

# Train the model
num_epochs = 10
batch_size = 64
optimizer = torch.optim.Adam(model.parameters(), lr=0.001) # Assuming you use Adam optimizer
criterion = torch.nn.CrossEntropyLoss() # Assuming a classification task

for epoch in range(num_epochs):
    model.train()
    epoch_loss = 0 # To track the total loss for the epoch
    for i in range(0, X_train_tensor.shape[0], batch_size):
        # Get the mini-batch
        batch_X = X_train_tensor[i:i+batch_size]
        batch_y = y_train_tensor[i:i+batch_size]

        # Zero the gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(batch_X)

        # Compute the Loss
        loss = criterion(outputs, batch_y)
        epoch_loss += loss.item() # Accumulate Loss for this epoch

    # Backward pass
    loss.backward()

    # Update parameters
    optimizer.step()

    # Append average epoch loss to the list
    avg_epoch_loss = epoch_loss / (X_train_tensor.shape[0] // batch_size)
    train_losses.append(avg_epoch_loss)
    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_epoch_loss:.4f}")

```

```

Epoch [1/10], Loss: 0.6738
Epoch [2/10], Loss: 0.6199
Epoch [3/10], Loss: 0.6197
Epoch [4/10], Loss: 0.6197
Epoch [5/10], Loss: 0.6196
Epoch [6/10], Loss: 0.6183
Epoch [7/10], Loss: 0.6135
Epoch [8/10], Loss: 0.6066
Epoch [9/10], Loss: 0.5993
Epoch [10/10], Loss: 0.5934


```

```

In [16]: # Plotting the training Loss
plt.figure(figsize=(10, 6))
plt.plot(range(1, num_epochs + 1), train_losses, marker='o', color='b', label='Training Loss')

```

```
plt.title("Training Loss Over Epochs")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



```
In [17]: # Evaluate on the test set
model.eval()
with torch.no_grad():
    test_outputs = model(X_test_tensor)
    _, predicted = torch.max(test_outputs, 1)
    accuracy = accuracy_score(y_test_tensor.cpu(), predicted.cpu())

print(f"Test Accuracy: {accuracy:.4f}")

# Print classification report
print(classification_report(y_test_tensor.cpu(), predicted.cpu(), target_names=['negative', 'neutral', 'positive']))
```

```
Test Accuracy: 0.9317
```

	precision	recall	f1-score	support
negative	0.45	0.27	0.33	162
neutral	0.30	0.11	0.16	300
positive	0.95	0.99	0.97	6464
accuracy			0.93	6926
macro avg	0.57	0.45	0.49	6926
weighted avg	0.91	0.93	0.92	6926

```
In [18]: # Function to predict sentiment for a real review
def predict_sentiment(review, model, vectorizer):
    # Preprocess the review
    review = preprocess_text(review)
    review = tokenize_and_remove_stopwords(review)

    # Vectorize the review
    review_vector = vectorizer.transform([review]).toarray()

    # Convert the review vector to tensor
    review_tensor = torch.tensor(review_vector, dtype=torch.float32)

    # Make prediction
    model.eval()
    with torch.no_grad():
        output = model(review_tensor)
        _, predicted = torch.max(output, 1)

    return ['negative', 'neutral', 'positive'][predicted.item()]

# Test with a real review
real_review = "I absolutely love this product! It works great and exceeded my expectations."
sentiment = predict_sentiment(real_review, model, vectorizer)
print(f"Predicted sentiment for the review: {sentiment}")
```

Predicted sentiment for the review: positive

```
In [19]: import gradio as gr

# Function to predict sentiment for a real review
def predict_sentiment(review, model, vectorizer):
    # Preprocess the review
    review = preprocess_text(review)
    review = tokenize_and_remove_stopwords(review)

    # Vectorize the review
    review_vector = vectorizer.transform([review]).toarray()
```



```

# Convert the review vector to tensor
review_tensor = torch.tensor(review_vector, dtype=torch.float32)

# Make prediction
model.eval()
with torch.no_grad():
    output = model(review_tensor)
    _, predicted = torch.max(output, 1)

    return ['negative', 'neutral', 'positive'][predicted.item()]

# Create Gradio interface
def interface_fn(review):
    sentiment = predict_sentiment(review, model, vectorizer)
    return f"Predicted sentiment for the review: {sentiment}"

# Define the Gradio interface
iface = gr.Interface(
    fn=interface_fn,
    inputs=gr.Textbox(label="Enter your review here"),
    outputs=gr.Textbox(label="Predicted Sentiment"),
    title="Amazon Product Review Sentiment Analysis",
    description="Enter an Amazon product review and get the predicted sentiment (negative, neutral, or positive).")

# Launch the interface
iface.launch()

```

Running on local URL: <http://127.0.0.1:7870>

To create a public link, set `share=True` in `launch()`.



Out[19]:

IMPORTANT: You are using gradio version 4.26.0, however version 4.44.1 is available, please upgrade.

Model 2; LSTM (Long Short-Term Memory) for Sentimental Analysis

In [20]:

```

# Tokenize reviews again for LSTM input
tokenized_reviews = [nltk.word_tokenize(review) for review in df['processed_reviews']]

# Build a vocabulary and encode the reviews as sequences of integers
all_words = [word for review in tokenized_reviews for word in review]
vocab = list(set(all_words))
vocab_size = len(vocab)

word_to_idx = {word: idx + 1 for idx, word in enumerate(vocab)} # Start indexing from 1
word_to_idx['<PAD>'] = 0 # Padding token

# Convert reviews to sequences of integers
def text_to_sequence(text, word_to_idx):
    return [word_to_idx.get(word, word_to_idx['<PAD>']) for word in text]

X_sequences = [text_to_sequence(review, word_to_idx) for review in tokenized_reviews]

# Pad the sequences to have uniform length (e.g., 100 words per review)
max_len = 100
X_padded = np.array([seq[:max_len] + [0] * (max_len - len(seq)) if len(seq) < max_len else seq[:max_len] for seq in X_sequences])

# Convert Labels to tensor
y = torch.tensor(df['sentiment'].map({'negative': 0, 'neutral': 1, 'positive': 2}).values, dtype=torch.long)

# Split the data into train and test sets (matching the previous train-test split)
X_train_tensor, X_test_tensor = torch.tensor(X_padded[:len(X_train)], dtype=torch.long), torch.tensor(X_padded[len(X_train):], dtype=torch.long)
y_train_tensor, y_test_tensor = y[:len(X_train)], y[len(X_train):]

# Create DataLoader for batching
train_data = TensorDataset(X_train_tensor, y_train_tensor)

```

```
test_data = TensorDataset(X_test_tensor, y_test_tensor)

train_loader = DataLoader(train_data, batch_size=64, shuffle=True)
test_loader = DataLoader(test_data, batch_size=64, shuffle=False)
```

```
In [21]: class LSTMModel(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim, n_layers, dropout_prob):
        super(LSTMModel, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers, batch_first=True, dropout=dropout_prob)
        self.fc = nn.Linear(hidden_dim, output_dim)
        # Dropout layer to prevent overfitting
        self.dropout = nn.Dropout(dropout_prob)

    def forward(self, x):
        embedded = self.embedding(x)
        lstm_out, (hidden, cell) = self.lstm(embedded)
        # We use the last hidden state as the representation for the entire sequence
        hidden = hidden[-1, :, :]
        out = self.fc(self.dropout(hidden))
        return out
```

```
In [22]: # Model hyperparameters
embedding_dim = 100
hidden_dim = 128
output_dim = 3 # 3 classes: negative, neutral, positive
n_layers = 2
dropout_prob = 0.5

# Instantiate the model
model = LSTMModel(vocab_size=len(word_to_idx), embedding_dim=embedding_dim, hidden_dim=hidden_dim,
                  output_dim=output_dim, n_layers=n_layers, dropout_prob=dropout_prob)

# Loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

# Train the model
num_epochs = 5
for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct_preds = 0
    total_preds = 0

    for batch in train_loader:
        # Get the batch data
        reviews, labels = batch
        optimizer.zero_grad()

        # Forward pass
        outputs = model(reviews)

        # Calculate Loss
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # Update Loss and accuracy metrics
        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct_preds += (predicted == labels).sum().item()
        total_preds += labels.size(0)

    epoch_loss = running_loss / len(train_loader)
    epoch_accuracy = correct_preds / total_preds
    print(f"Epoch {epoch + 1}/{num_epochs}, Loss: {epoch_loss:.4f}, Accuracy: {epoch_accuracy:.4f}")
```

```
Epoch 1/5, Loss: 0.3137, Accuracy: 0.9291
Epoch 2/5, Loss: 0.3045, Accuracy: 0.9300
Epoch 3/5, Loss: 0.3031, Accuracy: 0.9302
Epoch 4/5, Loss: 0.3010, Accuracy: 0.9303
Epoch 5/5, Loss: 0.3012, Accuracy: 0.9305
```

```
In [23]: # Evaluate the model on the test set
model.eval()
test_correct_preds = 0
test_total_preds = 0
test_running_loss = 0.0

with torch.no_grad():
    for batch in test_loader:
        reviews, labels = batch

        # Forward pass
        outputs = model(reviews)

        # Calculate Loss
        loss = criterion(outputs, labels)
        test_running_loss += loss.item()

        # Update accuracy metrics
        _, predicted = torch.max(outputs, 1)
        test_correct_preds += (predicted == labels).sum().item()
        test_total_preds += labels.size(0)

test_loss = test_running_loss / len(test_loader)
test_accuracy = test_correct_preds / test_total_preds
print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}")
```

Test Loss: 0.2654, Test Accuracy: 0.9454

```
In [24]: # Display classification report
from sklearn.metrics import classification_report
y_pred = []
y_true = []

with torch.no_grad():
    for batch in test_loader:
        reviews, labels = batch
        outputs = model(reviews)
        _, predicted = torch.max(outputs, 1)

        y_pred.extend(predicted.numpy())
        y_true.extend(labels.numpy())

print(classification_report(y_true, y_pred, target_names=['negative', 'neutral', 'positive']))
```

	precision	recall	f1-score	support
negative	0.00	0.00	0.00	149
neutral	0.10	0.00	0.01	222
positive	0.95	1.00	0.97	6555
accuracy			0.95	6926
macro avg	0.35	0.33	0.33	6926
weighted avg	0.90	0.95	0.92	6926

```
In [25]: # Test the model with a real review
def predict_sentiment_lstm(review, model, word_to_idx, max_len=100):
    # Preprocess the review
    review = preprocess_text(review)
    review = tokenize_and_remove_stopwords(review)

    # Convert review to sequence of integers
    review_sequence = text_to_sequence(review, word_to_idx)
    review_sequence = review_sequence[:max_len] + [0] * (max_len - len(review_sequence)) # Padding

    # Convert to tensor
    review_tensor = torch.tensor([review_sequence], dtype=torch.long)

    # Get the prediction
    model.eval()
    with torch.no_grad():
        output = model(review_tensor)
        _, predicted = torch.max(output, 1)

    return ['negative', 'neutral', 'positive'][predicted.item()]

# Test with a real review
real_review = "I love this product! It is fantastic and does exactly what I need."
predicted_sentiment = predict_sentiment_lstm(real_review, model, word_to_idx)
print(f'Predicted sentiment for the review: {predicted_sentiment}')
```

Predicted sentiment for the review: positive

```
In [21]: import gradio as gr
import torch

# Define the function for sentiment prediction using the LSTM model
def predict_sentiment_lstm(review, model, word_to_idx, max_len=100):
    # Preprocess the review (tokenization, stopword removal, etc.)
    review = preprocess_text(review)
    review = tokenize_and_remove_stopwords(review)

    # Convert review to sequence of integers
    review_sequence = text_to_sequence(review, word_to_idx)
    review_sequence = review_sequence[:max_len] + [0] * (max_len - len(review_sequence)) # Padding

    # Convert to tensor
    review_tensor = torch.tensor([review_sequence], dtype=torch.long)

    # Get the prediction
    model.eval()
    with torch.no_grad():
        output = model(review_tensor)
        _, predicted = torch.max(output, 1)

    return ['negative', 'neutral', 'positive'][predicted.item()]

# Gradio interface function to predict sentiment
def gradio_predict(review):
    return predict_sentiment_lstm(review, model, word_to_idx)

# Set up the Gradio interface
iface = gr.Interface(
    fn=gradio_predict,
    inputs="text",
    outputs="text",
    title="Sentiment Prediction using LSTM",
    description="Enter a review to predict its sentiment (negative, neutral, or positive).")

# Launch the Gradio app
iface.launch()
```

Running on local URL: <http://127.0.0.1:7872>
IMPORTANT: You are using gradio version 4.26.0, however version 4.44.1 is available, please upgrade.

To create a public link, set `share=True` in `launch()`.



Out[21]:

In []:

Model 3: Pretrained DistilBERT Model for Sentiment Analysis

In []:

```
In [8]: import gradio as gr
import torch
from transformers import DistilBertTokenizer, DistilBertForSequenceClassification

# Load the pretrained DistilBERT model and tokenizer
model = DistilBertForSequenceClassification.from_pretrained('distilbert-base-uncased', num_labels=3)
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)

# Define the sentiment prediction function
def predict_sentiment(review):
    model.eval()
    # Tokenize the review
    encodings = tokenizer(review, padding=True, truncation=True, max_length=100, return_tensors="pt")

    # Move inputs to the same device as the model
    input_ids = encodings['input_ids'].to(device)
    attention_mask = encodings['attention_mask'].to(device)

    # Get the model's prediction
    with torch.no_grad():
        outputs = model(input_ids, attention_mask=attention_mask)
        logits = outputs.logits
        _, predicted = torch.max(logits, 1)

    sentiment = ['negative', 'neutral', 'positive'][predicted.item()]
    return sentiment

# Create Gradio interface
iface = gr.Interface(fn=predict_sentiment,
                    inputs=gr.Textbox(label="Enter Review", placeholder="Type your review here...", lines=2),
                    outputs=gr.Textbox(label="Predicted Sentiment"),
                    title="Sentiment Analysis of Product Reviews",
                    description="Enter a product review, and the model will predict whether the sentiment is 'negative', 'neutral', or 'positive'",
                    live=True)

# Launch the Gradio interface
iface.launch()
```

In []:

Model 4; Pretrained BERT Model for sentimental Analysis

```
In [28]: # Preprocessing your data
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

# Tokenizing the text data
def tokenize_reviews(reviews, tokenizer, max_len=100):
    return tokenizer(reviews.tolist(), padding=True, truncation=True, max_length=max_len, return_tensors="pt")
```

```

# Convert sentiment labels to numeric values
y = df['sentiment'].map({'negative': 0, 'neutral': 1, 'positive': 2}).values

# Split data into train and test
X_train, X_test, y_train, y_test = train_test_split(df['processed_reviews'], y, test_size=0.2, stratify=y, random_state=42)

# Tokenize the reviews
train_encodings = tokenize_reviews(X_train, tokenizer)
test_encodings = tokenize_reviews(X_test, tokenizer)

# Convert labels to tensors
train_labels = torch.tensor(y_train)
test_labels = torch.tensor(y_test)

```

In [29]: `from torch.utils.data import DataLoader, TensorDataset`

```

# Create DataLoader for training and testing
train_data = TensorDataset(train_encodings['input_ids'], train_encodings['attention_mask'], train_labels)
test_data = TensorDataset(test_encodings['input_ids'], test_encodings['attention_mask'], test_labels)

train_loader = DataLoader(train_data, batch_size=32, shuffle=True)
test_loader = DataLoader(test_data, batch_size=32, shuffle=False)

```

In [30]: `from transformers import BertForSequenceClassification, AdamW`

```

# Load pretrained BERT model for sequence classification
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=3)

# Define optimizer and loss function
optimizer = AdamW(model.parameters(), lr=1e-5)

# Move model to GPU if available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```

Out[30]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
  (pooler): BertPooler(
    (dense): Linear(in_features=768, out_features=768, bias=True)
    (activation): Tanh()
  )
  (dropout): Dropout(p=0.1, inplace=False)
  (classifier): Linear(in_features=768, out_features=3, bias=True)
)

```

In [9]: `from torch import nn`
`import torch.optim as optim`

```

# Train the model
def train_model(model, train_loader, optimizer, device, num_epochs=3):
    model.train()
    for epoch in range(num_epochs):
        running_loss = 0.0
        correct_preds = 0
        total_preds = 0

        for batch in train_loader:
            # Get the inputs and labels from the batch
            input_ids, attention_mask, labels = batch
            input_ids, attention_mask, labels = input_ids.to(device), attention_mask.to(device), labels.to(device)

```

```

        # Zero gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(input_ids, attention_mask=attention_mask, labels=labels)
        loss = outputs.loss
        logits = outputs.logits

        # Backward pass
        loss.backward()
        optimizer.step()

        # Track Loss and accuracy
        running_loss += loss.item()
        _, predicted = torch.max(logits, 1)
        correct_preds += (predicted == labels).sum().item()
        total_preds += labels.size(0)

    epoch_loss = running_loss / len(train_loader)
    epoch_accuracy = correct_preds / total_preds
    print(f"Epoch {epoch+1}/{num_epochs}, Loss: {epoch_loss:.4f}, Accuracy: {epoch_accuracy:.4f}")

# Train the model for a few epochs
train_model(model, train_loader, optimizer, device, num_epochs=3)

```

```

In [10]: from sklearn.metrics import classification_report

# Evaluate the model
def evaluate_model(model, test_loader, device):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for batch in test_loader:
            input_ids, attention_mask, labels = batch
            input_ids, attention_mask, labels = input_ids.to(device), attention_mask.to(device), labels.to(device)

            # Forward pass
            outputs = model(input_ids, attention_mask=attention_mask)
            logits = outputs.logits

            # Get predictions
            _, predicted = torch.max(logits, 1)
            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    # Print classification report
    print(classification_report(all_labels, all_preds, target_names=['negative', 'neutral', 'positive']))

# Evaluate the model
evaluate_model(model, test_loader, device)

```

```

In [11]: def predict_sentiment(review, model, tokenizer, device, max_len=100):
    model.eval()
    # Tokenize the review
    encodings = tokenizer(review, padding=True, truncation=True, max_length=max_len, return_tensors="pt")

    # Move inputs to the same device as the model
    input_ids = encodings['input_ids'].to(device)
    attention_mask = encodings['attention_mask'].to(device)

    # Get the model's prediction
    with torch.no_grad():
        outputs = model(input_ids, attention_mask=attention_mask)
        logits = outputs.logits
        _, predicted = torch.max(logits, 1)

    return ['negative', 'neutral', 'positive'][predicted.item()]

# Test the model with a real review
real_review = "I absolutely love this product! It works as expected and more."
predicted_sentiment = predict_sentiment(real_review, model, tokenizer, device)
print(f"Predicted sentiment: {predicted_sentiment}")

Predicted sentiment: neutral

```

```

In [23]: import gradio as gr
import torch
from transformers import BertTokenizer, BertForSequenceClassification

# Load the pretrained BERT model and tokenizer
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=3)
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

# Move model to device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)

# Define the sentiment prediction function
def predict_sentiment(review):
    model.eval()
    # Tokenize the review
    encodings = tokenizer(review, padding=True, truncation=True, max_length=100, return_tensors="pt")

    # Move inputs to the same device as the model
    input_ids = encodings['input_ids'].to(device)
    attention_mask = encodings['attention_mask'].to(device)

```

```

# Get the model's prediction
with torch.no_grad():
    outputs = model(input_ids, attention_mask=attention_mask)
    logits = outputs.logits
    _, predicted = torch.max(logits, 1)

    sentiment = ['negative', 'neutral', 'positive'][predicted.item()]
    return sentiment

# Create Gradio interface
iface = gr.Interface(fn=predict_sentiment,
                    inputs=gr.Textbox(label="Enter Review", placeholder="Type your review here...", lines=2),
                    outputs=gr.Textbox(label="Predicted Sentiment"),
                    title="Sentiment Analysis of Product Reviews",
                    description="Enter a product review, and the model will predict whether the sentiment is 'negative', 'neutral', or 'positive'",
                    live=True)

# Launch the Gradio interface
iface.launch()

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight']
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
 Running on local URL: http://127.0.0.1:7874

To create a public link, set `share=True` in `launch()`.



Out[23]:

IMPORTANT: You are using gradio version 4.26.0, however version 4.44.1 is available, please upgrade.

Model Analysis

In [6]: `import matplotlib.pyplot as plt`

```

# Model names and accuracies
models = ['FFNN', 'LSTM', 'BERT']
accuracies = [93.17, 95, 96]


# Create bar graph
plt.figure(figsize=(8, 6))
plt.bar(models, accuracies, color=['blue', 'green', 'orange'])

# Add titles and labels
plt.title('Model Accuracy Comparison', fontsize=16)
plt.xlabel('Models', fontsize=14)
plt.ylabel('Accuracy (%)', fontsize=14)

# Display accuracy values on top of bars
for i, accuracy in enumerate(accuracies):
    plt.text(i, accuracy + 0.5, f'{accuracy}%', ha='center', fontsize=12)

# Show plot
plt.tight_layout()
plt.show()

```



In []: